ReflectScan

A smartphone-based data-collection and RL-estimation prototype for the national highway network.

Any patrol vehicle + any Android phone = scalable data capture for RL assessment.

A proof-of-concept platform that aims to complement ₹80 Lakh – ₹1.5 Crore specialized retroreflectivity vehicles with a smartphone workflow mounted on existing patrol fleets. Physics-derived approach (distinct from pure ML regression). Offline-first data pipeline. End-to-end prototype exercised on local roads in Ahmedabad on 22 April 2026 with the phone secured on a two-wheeler mobile-phone stand, for pipeline validation; four-wheeler patrol-vehicle testing at highway speed and RL calibration against IRC-standard instruments are scheduled as Phase 1 next steps.

PoC	51	~₹1.5 K	Android 8+
CURRENT MATURITY — PROOF OF CONCEPT	AUTOMATED CODE-LEVEL TESTS PASSING	PROJECTED MARGINAL UNIT COST (MOUNT ONLY)	MINIMUM TARGET PLATFORM

What this submission contains

- **Prototype software:** Android app (Flutter) + FastAPI + PostGIS backend, end-to-end data pipeline working offline and online.
- **Pipeline-validation ride:** 82 GPS-tagged captures collected with the phone secured on a two-wheeler handlebar mobile-phone stand in Ahmedabad, exercising the camera → flash → classification → offline storage path on real road surfaces.
- **Physics-based RL formulation** documented step-by-step; correctness depends on subsequent calibration work.
- **Validation roadmap:** controlled testing plan against a handheld retroreflectometer (IRC 67 / IRC 35), not yet executed.
- **Public GitHub repository** with full source, 51 passing tests, and deployment documentation.

MATURITY DISCLAIMER

ReflectScan is a hackathon prototype. This document describes the system design, an end-to-end software pipeline, and a pipeline-validation ride (two-wheeler with mobile-phone stand). **It does not claim accuracy-validated or IRC-certified RL measurements.** Calibration against standard instruments is the first item in the proposed NHAI pilot — see pages 7 and 10.

Team Leader: Divyang Kuntar
Email: divyangkuntar@gmail.com
Mobile: 7016509214
Team size: 2
Repository: https://github.com/kuntardivyang/reflectivity
Live dashboard: https://reflectivity-3x1c.vercel.app
Screenshots, videos & APK: https://drive.google.com/drive/folders/1xpssEgShkJ3bRBdnN1gJgEUiT1i-ei5d

CONTEXT

The Problem NHAI Asked Us to Address

NHAI operates **1.46 Lakh km of national highways**. Each kilometre contains retroreflective assets — pavement markings, gantry signs, delineators, road studs — that must meet IRC 67 / IRC 35 minimum RL thresholds throughout their service life. When those thresholds fall, night-time visibility drops and accident risk rises.

Why current measurement approaches face scale limits

Approach	Typical unit cost	Practical annual coverage	Operator considerations
Handheld retroreflectometer (e.g. Delta-E LTL, Zehntner RETRORED)	₹2–5 Lakh per unit	Limited by staff-hours on the road	RISK inspector on live highway
Vehicle-mounted mobile retroreflectometer	₹80 Lakh – ₹1.5 Crore per vehicle	Limited by fleet size and dedicated crews	CAPEX specialist vehicle + crew
Pure ML / camera-only estimation (academic literature)	Low hardware cost	Theoretically unlimited	TRUST unauditable regression

Observed outcome: because each approach is either costly or slow, a significant share of the NH network is not measured on a regular cadence. A smartphone-based data-capture workflow could complement existing methods — *not replace certified instruments* — by broadening the sampling base at low marginal cost.

What the NHAI problem statement asks for

1. Retroreflectivity measurement by **vehicle / drone-mounted equipment**, to improve speed of coverage.
2. RL determination using **AI & ML techniques** on camera data.
3. A combination of the above.

ReflectScan addresses all three directions while explicitly acknowledging:

- This is a **prototype** designed to prove the data-capture and computation pipeline;
- Accuracy of RL output **depends on calibration**, which is a scheduled next step;
- The system is intended to **augment** IRC-compliant instruments, not substitute them in their certified role.

Target operating conditions and honest behaviour in each

The NHAI problem statement lists day/night, dry/wet, with/without street light, and fog. The prototype's intended behaviour in each is:

Condition	Design intent — prototype behaviour
Night + flash (primary operating mode)	Flash-delta RL computed; output subject to calibration.
Daylight	UNCAL tag emitted; GPS trace recorded; RL value not trusted by the dashboard.
Rain / wet asphalt	Design provides for a wet-geometry correction factor; empirical validation pending.
Dense fog	Flash signal attenuated; prototype's confidence reporting under fog is a subject of future testing.

SOLUTION

System Architecture

Three layers, each independently deployable. The mobile app functions offline; the backend and dashboard ingest and visualise what the fleet captures.

```
┌──────────────────────────────────────────────────────────────────┐
│  LAYER 1 — PATROL VEHICLE  (existing NHAI fleet, no new hardware)  │
│                                                                    │
│   Android phone on windshield mount                                │
│   ├─ Camera YUV420 stream @ 30 fps                                 │
│   ├─ LED flash strobe @ 2 Hz via CameraController.setFlashMode     │
│   ├─ GPS (GNSS + WiFi + cell hybrid, 10 s first-fix target)        │
│   │                                                                │
│   ├─►  YOLOv8n TFLite  →  lane / edge / crossing bounding boxes    │
│   ├─►  Luminance Δ  =  illuminated frame − ambient frame (Y-plane) │
│   ├─►  OECF curve  →  pixel value to cd/m²   (generic curve by default │
│   │                                            — per-device calibration│
│   │                                            required for accuracy)  │
│   ├─►  RL  =  ΔL / E(flash, distance) × 1000   mcd/m²/lux          │
│   ├─►  Physical clamp [0, 2000]; UNCAL tag if flash delta invalid  │
│   ├─►  Preliminary classification:                                 │
│   │       SAFE > 100 / WARNING 54–100 / CRITICAL < 54  (IRC-aligned)  │
│   └─►  Persist to on-device SQLite  →  upload queue                │
└──────────────────────────────────────────────────────────────────┘
                              │  batch sync when connectivity returns
                              ▼
┌──────────────────────────────────────────────────────────────────┐
│  LAYER 2 — CLOUD BACKEND  (FastAPI + PostgreSQL 15 + PostGIS 3.4)  │
│                                                                    │
│   POST /api/measurements   →  validate, aggregate to 100 m segments │
│   GET  /api/segments       →  GeoJSON FeatureCollection (map)      │
│   GET  /api/alerts         →  active WARNING / CRITICAL list       │
│   GET  /api/stats          →  KPI summary (per state, per NH, network) │
│                                                                    │
│   Alert engine:  on CRITICAL segment, raise ticket; skip duplicates │
│   UNCAL filter:  uncalibrated points excluded from rl_avg & status  │
└──────────────────────────────────────────────────────────────────┘
                              │
                              ▼
┌──────────────────────────────────────────────────────────────────┐
│  LAYER 3 — DASHBOARD  (Next.js + Leaflet + react-leaflet)          │
│                                                                    │
│   Live map:  green / yellow / red / grey highway segments          │
│   Alert feed:  chronological, filterable by division               │
│   KPI panel:  % SAFE / WARN / CRIT, coverage %, points this week   │
│   Export:  CSV (raw), PDF (executive summary per state)            │
└──────────────────────────────────────────────────────────────────┘
```

Data flow for a single capture

1. Camera delivers a YUV420 frame while the LED flash is ON (illuminated).
2. Next frame is captured with flash OFF (ambient baseline).
3. YOLOv8n returns a bounding box around the road marking. (If model asset is absent, a brightest-patch heuristic is used.)
4. Y-plane pixels inside the box are averaged for both frames.
5. An OECF lookup table converts pixel values to cd/m². The prototype ships with a generic gamma-2.2 curve; per-device calibration is pending.
6. Δ Luminance = illuminated − ambient.
7. Flash illuminance at marking = `flash_lumens / (4π × distance²)`. Uses a default flash lumen value; per-device photometric calibration is pending.
8. RL = (ΔL ÷ illuminance) × 1000. Clamped to [0, 2000].
9. GPS-tag, classify (with UNCAL fallback), persist to SQLite, queue for upload.

ACCURACY CAVEAT

Every step above is deterministic and reproducible from source. The *numerical RL output*, however, depends on two device-specific reference values — the per-phone OECF curve and the calibrated flash lumen output — that the current prototype uses generic defaults for. Establishing accuracy therefore requires a controlled calibration campaign, which is described in the "Validation Plan" section of this document.

METHODOLOGY

How ReflectScan Attempts to Estimate Retroreflectivity

Retroreflectivity (R_L) is defined in ASTM E1710 as the coefficient of retroreflected luminance — luminance of the marking in cd/m² divided by the illuminance reaching the marking in lux, scaled to mcd/m²/lux. ReflectScan's intent is to compute an *estimate* of this quantity on commodity smartphones using the same optical relationship that a certified retroreflectometer uses.

The optical chain

1. **Illumination:** phone LED delivers lumen output (nominal value; per-phone calibration required).
2. **Target:** road marking glass beads retroreflect the light back toward the phone.
3. **Sensor:** camera's Y-plane captures luminance-proportional pixel values.
4. **OECF curve:** maps 0–255 pixel values back to photometric cd/m². Generic gamma curve today; per-device table needed for accuracy.
5. **Ambient subtraction:** flash-OFF frame is subtracted to isolate the retroreflected component.

Core formula

```
ΔL   = L_illuminated − L_ambient
E    = flash_lumens / (4 π d²)
R_L  = (ΔL / E) × 1000
```

Units: L in cd/m², E in lux, R_L in mcd/m²/lux.

Clamp: output bounded to [0, 2000]. Real thermoplastic paint tops out near 500; the 2000 ceiling catches sensor saturation without producing runaway values.

Why we prefer a physics-derived approach over pure ML regression

Pure ML / CNN regression	ReflectScan (physics + ML detector)
Learns a function *photo → RL number*; weights are opaque.	Applies ASTM E1710 relationship directly; every step visible and reproducible.
A single unrepresentative training example can bias all outputs.	Calibration curves are derived, once, from known physical targets under controlled conditions.
No natural mechanism to decline a reading; always outputs a number.	Refuses to emit a classification when conditions are invalid (UNCAL tag).
Transfer to new paint formulations / new cameras typically requires retraining.	New phone model → a repeatable calibration procedure; core algorithm unchanged.

What ML does in our design

The YOLOv8n TFLite model's role is **detection, not RL estimation** — it locates the marking in the frame. Separation of concerns: ML locates where to sample; physics (subject to calibration) computes the number from those samples. If the model asset is absent, a brightest-patch heuristic takes over — under flash conditions, retroreflective paint is typically much brighter than surrounding asphalt, so the brightest patch in the lower frame tends to coincide with the marking. Robustness of this heuristic across lighting conditions is itself subject to validation.

> **Design position:** ReflectScan is architected so that the RL pipeline is auditable end-to-end. It does *not* claim current numerical parity with IRC-certified instruments. Calibrating the pipeline to achieve that parity is the primary work item in the proposed pilot.

INNOVATION

What Distinguishes ReflectScan

> **Not another ML black box.** Every competing smartphone-based approach in the academic literature trains a neural network to guess RL from a photo. ReflectScan instead *computes* RL using the same photometric physics a certified retroreflectometer uses — on hardware NHAI already owns, at a marginal cost of ~₹1 500 per vehicle.

Four design contributions

1. Physics-derived RL pipeline, not ML-regressed RL

Academic proposals frequently train a CNN to output RL directly from an image. This is a black box and resists audit. ReflectScan instead *computes* RL using the same physical relationship a certified retroreflectometer uses — luminance delta, OECF conversion, inverse-square-law illuminance. The result is a reproducible, explainable number whose *correctness under calibration* an optical engineer can examine.

2. Explicit uncertainty handling — the UNCAL tag

When flash differential is unavailable (strong daylight, torch hardware failure), the system **refuses** to output a SAFE / WARNING / CRITICAL classification and emits `UNCAL` instead. RL is clamped to the physically plausible ceiling [0, 2000] mcd/m²/lux; backend aggregation excludes UNCAL points from segment averages. Commercial retroreflectometers do not routinely expose their own uncertainty to the caller; ReflectScan does.

3. Zero-marginal-cost fleet deployment model

NHAI already operates thousands of patrol vehicles. The software runs on any Android 8+ phone with a flash. The projected marginal cost of deploying a new surveying unit is **~₹1 500** (a vibration-damped windshield mount). This is a design claim about deployability; actual field durability of the mount and smartphone workflow remains to be observed at pilot scale.

4. Offline-first data pipeline

Rural NH sections have intermittent coverage. Every capture writes to on-device SQLite first; uploads occur in batches when connectivity returns. This behaviour was exercised in the 22 April 2026 prototype drive and functions as designed in our code-level tests.

Head-to-head positioning

	Delta-E LTL-XL handheld	Zehntner RETRORED vehicle	Pure ML research	ReflectScan (this prototype)
Typical unit cost	₹3 Lakh	₹1 Cr +	Low	**~₹1 500 (projected)**
Physics-derived	Yes	Yes	No	**Yes**
Auditable formula	Yes	Yes	No	**Yes**
Accuracy validation	IRC-certified	IRC-certified	Benchmark datasets	**Planned (Phase 1)**
Vehicle-mounted workflow	No	Yes (dedicated)	Possibly	**Yes (any vehicle, by design)**
Offline operation	Yes	Yes	No	**Yes**
Uncertainty flagging	No	No	No	**Yes (UNCAL)**
Indian-road model tuning	N/A	N/A	Rare	**Planned (IDD + Roboflow)**

Innovation summary: ReflectScan's contribution is the *combination* of physics-derived pipeline explicit uncertainty low-marginal-cost deployment model offline-first data path in a single Android-phone artefact that runs on hardware NHAI already owns — as a complement to, not a replacement for, IRC-certified instruments.

FEASIBILITY

What Is Built, and the Limits of What Has Been Tested

This is a working prototype, not a certified product. Below: what is complete, what is in progress, and what remains to be validated.

Component status

Component	Status	Evidence
Backend (FastAPI + PostGIS)	CODE COMPLETE	27 pytest tests passing; 5 REST endpoints live; OpenAPI spec auto-generated
Mobile app (Flutter + Riverpod)	CODE COMPLETE	24 unit tests passing; APK installed and exercised on a real Android device
End-to-end pipeline test ride	PIPELINE EXERCISED	82 GPS-tagged captures from a two-wheeler with mobile-phone stand, 22 April 2026, Ahmedabad local roads (pipeline validation only; see next section)
RL accuracy vs IRC instrument	NOT YET ASSESSED	Planned as Phase 1 of the pilot; requires access to a certified retroreflectometer
AI model (YOLOv8n TFLite, Indian-road tuned)	IN PROGRESS	Stock COCO weights loaded as baseline; IDD + Roboflow fine-tune planned
Dashboard (Next.js + Leaflet)	CODE COMPLETE	Next.js + Leaflet dashboard built and serving; map, KPI panel, and alert feed operational (HTTP 200 verified, 23 April 2026)
Documentation	COMPLETE	README, architecture.md, research_findings.md, problem_statement.md

Already deployed — not a simulation. The backend API is live at `https://reflectivity-production.up.railway.app` (Railway, PostGIS, seeded with Ahmedabad NH-48 data). The dashboard is publicly accessible at `https://reflectivity-3x1c.vercel.app`. A judge can open both URLs right now. A pilot can begin the moment calibration data is available — no infrastructure setup is required.

Deployment roadmap

Phase	Timeline	Scope	Projected cost
Phase 0 — Calibration & IRC benchmark	Month 1	Controlled testing on reference markings with a handheld retroreflectometer; per-phone OECF + flash calibration; error characterisation	₹50 K instrument-hire + field-test logistics
Phase 1 — Single-state pilot	Month 2–3	~200 patrol phones, one state (e.g. Gujarat or Maharashtra), ~5 000 km NH coverage attempt, ground-truth subset validated against an IRC handheld device	₹3 Lakh one-time + ~₹40 K/month ops (estimate)
Phase 2 — RAMS integration	Month 4	Integration with NHAI Road Asset Management System; maintenance-ticket workflow from CRITICAL alerts; user training for 10 regional engineers	₹2 Lakh dev + ongoing ops (estimate)
Phase 3 — Broader rollout	Month 5+	Progressive rollout subject to pilot outcomes and NHAI approval	To be costed after pilot

Hardware requirements — minimal

Item	Spec	Indicative cost
Android phone	Android 8+, rear camera with LED flash	Reuse fleet inventory where possible
Windshield mount	Vibration-damped, suction / adhesive	~₹1 500 per vehicle
Backend VPS	4 vCPU / 8 GB / 100 GB SSD + Postgres 15 + PostGIS 3.4	~₹4 000 / month
Dashboard hosting	Static Next.js on a CDN free tier	~₹0

Known challenges and how the prototype addresses them today

Challenge	Prototype approach
Per-phone camera variance	OECF lookup table intended to be calibrated per device model (pending); prototype ships a generic gamma-2.2 curve as a placeholder
Daylight drowns flash signal	UNCAL tag; day-time captures preserve GPS path but are excluded from RL statistics
Torch API conflicts with active camera stream	Flash routed through `CameraController.setFlashMode`, which shares the camera session; auto-disables after repeated torch errors
GPS weak on cloudy starts	Hybrid mode (GNSS + Wi-Fi + cell) plus one-shot fast-fix in parallel with the stream
Offline operation	SQLite-backed upload queue with exponential-backoff retry
Implausible RL output on invalid input	Physical clamp [0, 2000 mcd/m²/lux] inside the calculator itself

PROTOTYPE VALIDATION SCOPE

What Works Today — End-to-End Pipeline Verified

> ReflectScan is a working, installed, exercised prototype — not a concept or a simulation. Every software component is code-complete, deployed, and independently verifiable via the live links on page 12.

What has been tested

- **Code-level tests.** 27 backend pytest tests and 24 mobile Dart unit tests execute on every build; all passing at time of submission.
- **Offline data-capture pipeline.** On 22 April 2026, the app was installed on a physical Android device, secured on a two-wheeler handlebar mobile-phone stand, and exercised on *local road conditions in Ahmedabad* for approximately 14 minutes. The exercise produced **82 GPS-tagged captures** that were written to the on-device SQLite store, confirming the camera, flash control, detection fallback, classification, and offline-persistence paths execute end-to-end. A four-wheeler / dashboard-mount drive at patrol cruise speed is the next test condition (see Validation Plan, page 10).
- **Error and edge-case handling.** Torch failure breaker, GPS-delayed start, permission-denied flows, and missing-model fallback (brightest-patch heuristic) were all observed to behave as designed during the two-wheeler prototype ride.

What has *not* been tested

- **Measurement accuracy against an IRC-certified retroreflectometer.** No comparison reading from a Delta-E LTL-XL, Zehntner RETRORED, or equivalent has been collected. *Any RL number produced by the current prototype is therefore indicative, not authoritative.*
- **Operation on NHAI-owned national highways.** The 22 April drive was on *local Ahmedabad roads*, which contain comparable pavement markings but were not operated on under NHAI supervision, and do not constitute a national-highway pilot.
- **Per-device OECF calibration.** The prototype uses a generic gamma-2.2 curve. A proper calibration requires controlled-illumination reference targets.
- **Photometric flash calibration.** The flash lumen value is a nominal default for mid-range Android hardware; it has not been measured per device.
- **Wet, foggy, and low-light conditions.** Design intent is documented; empirical behaviour across these conditions remains to be measured.
- **Mount durability at highway cruise speeds.** Prototype drive was at urban speeds; high-speed vibration / wind-load effects on capture quality are not yet characterised.
- **Multi-device consistency.** All captures were collected on one device model. Inter-device variance is unknown.

What the 82 captures *do* and *do not* demonstrate

The drive does demonstrate	The drive does not demonstrate
End-to-end software pipeline works on real hardware under real road conditions.	That the RL numbers produced are within any stated tolerance of a certified instrument.
Flash strobe, GPS fix, frame pairing, detection fallback, and SQLite persistence function together.	That the chosen SAFE/WARNING/CRITICAL thresholds classify this device's output correctly.
GPS accuracy sufficient for 100 m segment aggregation on local roads.	Behaviour on expressways, in rain, in fog, or at speeds > 50 km/h.
Offline queue + later sync works as designed.	Durability of the mount or camera-app combination under continuous 8-hour patrol use.

Bottom line. ReflectScan at submission time is an end-to-end working prototype whose software pipeline has been exercised on local Indian road conditions. It is *not* an accuracy-validated retroreflectivity measurement instrument. Establishing accuracy against IRC-compliant standards is the subject of the proposed Phase 0 calibration and Phase 1 pilot.

SCALABILITY & SUSTAINABILITY

Designed to Scale, Subject to Validation

> **Scaling ReflectScan costs almost nothing.** Each new surveying unit is one windshield mount (~₹1 500). No new vehicles. No new crews. No software licence fees. NHAI's existing patrol fleet becomes a dense retroreflectivity sensor network the moment the app is installed.

The cost and coverage projections below are **design estimates** derived from architecture capacity and unit-hardware pricing. They assume the prototype graduates to a calibrated tool through Phase 0 and Phase 1.

Scale model

Fleet size	Design coverage frequency	Projected annual cost	Relative data volume
1 phone (pilot)	Device-limited	~₹1 500 (mount)	1× baseline
200 phones (single state)	~1 Lakh km / year (design target)	~₹3 Lakh (estimate)	~200×
2 000 phones (national, subject to pilot outcomes)	Entire NH network on a multi-week cadence (design target)	~₹30 Lakh (estimate)	~2 000×

> **Scaling property:** each incremental phone adds an incremental windshield mount and no per-unit software cost. Cost scales *linearly* with fleet size while coverage compounds. These are projected figures; they require observed per-phone throughput from the pilot to finalise.

Projected 5-year cost (design estimates)

Comparison baseline: 50 specialized vehicles at ₹80 Lakh each plus ongoing crew, fuel, calibration. ReflectScan baseline: 2 000 phones × ₹1 500 mount + backend hosting + admin. Exact totals depend on NHAI existing-vehicle reuse, negotiated software-hosting rates, and pilot findings. Figures are illustrative, not quotations.

Architectural capacity (projections from design)

- **Edge:** measurement runs on-device; backend sees ~100 bytes per point. Hypothetical 2 000 phones × 500 points/day ≈ **100 MB/day** total ingress.
- **Cloud:** PostGIS with spatial partitioning; 14 crore points (India full coverage × 1 year) fits within a single managed Postgres instance. Not yet stress-tested; subject to load-testing during Phase 1.
- **Dashboard:** OpenStreetMap tile CDN serves the base map; the Next.js app is stateless and can be scaled horizontally behind a load balancer.

Potential data re-use

Use case	How ReflectScan would enable it
Predictive maintenance	Time-series on per-segment RL decay could flag segments approaching threshold; statistical modelling is a future work item.
Contractor-performance audit	Re-striping dates set segment birth times; degradation rate per contractor becomes measurable.
Warranty enforcement	CRITICAL findings in the warranty window can generate evidence packets (GPS, timestamp, photo) for chargebacks.
Public safety index	Per-district aggregates could be published as an open-data indicator, comparable to AQI in spirit.
Smart-highway integration	Alerts can feed NHAI HTMC / VMS systems in low-visibility zones.

Sustainability

Zero new hardware implies zero new e-waste at scale. Capture piggybacks on patrol drives NHAI is already performing, with no additional fuel burn. Open-source Flutter + Python stack avoids licence fees and vendor lock-in. Phones retire on standard NHAI replacement cycles and the software inherits the upgrade automatically.

PROTOTYPE DETAIL

Mobile App, AI/ML Component, and Vehicle Integration

Mobile app screens & user roles

Screen	Purpose	Key elements
Permission gate	First launch	Camera + location request with rationale + retry on denial
Live survey	Main workflow	Camera preview, strobing flash, detection overlay, status HUD (SAFE / WARNING / CRITICAL / UNCAL), GPS pill, captured / sent counters, START/STOP button, DEMO MODE banner when the ML model asset is absent
History	Past surveys	Session list (highway tag, duration, point count), drill-down to session detail
Error / retry	Failure states	Human-readable message, retry button — no silent failure

Role	Access
Patrol driver / surveyor	Start/stop survey, view history on own phone
Regional engineer	Read-only dashboard for their division; acknowledge alerts
NHAI admin	Full dashboard, export CSV/PDF, configure alert rules

AI / ML component

Model: YOLOv8n, TFLite float16, ~6 MB. Small enough to ship inside the APK; on typical mid-range Android hardware, inference is reported by the framework at interactive rates. Actual throughput on NHAI-issue devices is a future measurement.

Planned classes (kept minimal so the model can generalise from a modest dataset):

- `lane_line` — dashed centre / lane dividers
- `edge_line` — solid white highway edge
- `crossing` — zebra / pedestrian marking

Training data strategy (planned):

- **Base:** Indian Driving Dataset (IDD, IIIT-H) — detection-labelled Indian road imagery
- **Augment:** Roboflow Universe road-marking datasets — already in YOLO format
- **Augmentations:** mosaic, random brightness (day/night simulation), gaussian blur (motion blur simulation)
- **Target:** mAP@0.5 > 0.6 on held-out Indian highway images (a *design target*, to be verified post-training)
- **Recipe:** `yolo train model=yolov8n.pt data=roads.yaml epochs=100 imgsz=640` on a free-tier Colab T4

> **Graceful degradation:** if the TFLite asset is absent, the app runs a brightest-patch heuristic on a 5×3 grid across the lower 2/3 of the frame. Under flash, retroreflective paint is typically much brighter than asphalt; the brightest patch tends to coincide with the marking. This behaviour was observed during the two-wheeler prototype ride. Quantitative characterisation of its accuracy across lighting conditions remains to be done.

Vehicle integration

Parameter	Design value
Mount position	Centre dashboard, ~1.0–1.2 m above road, facing forward
Intended cruise speed	40 – 80 km/h (standard patrol speed)
Frame rate	30 fps YUV420 — camera API native, no conversion
Capture throttle	One measurement per 800 ms (SQLite-friendly)
Geometry	~5 m forward look in the current prototype; ASTM 30 m geometry is a future extension

Design considerations — speed vs capture density

Speed (design target)	Expected samples per km	Notes
30 km/h	~150	Urban / tolled stretches
60 km/h	~75	Typical patrol cruise
90 km/h	~50	Expressway cruise (untested)

Table describes expected capture density derived from the 800 ms throttle, not measured accuracy. Empirical noise characterisation per speed band is a deliverable of the Phase 1 pilot.

Drone extension (future mode)

The same pipeline can run on a drone pointing its camera at gantry signs and elevated structures that patrol cars cannot easily reach. Drone and vehicle feeds land in the same backend; the dashboard is modality-agnostic. Drone integration is future work; not part of the current submission.

Validation Plan & Next Steps

This section describes how ReflectScan would progress from a working prototype to a tool whose outputs can be trusted by NHAI engineers.

Phase 0 — Controlled calibration (Month 1)

1. **Procure a reference instrument** — a handheld retroreflectometer conforming to IRC 67 / IRC 35 / ASTM E1710 (e.g. Delta-E LTL-XL, Zehntner RETRORED). Equipment can typically be hired or borrowed from a state PWD / NHAI regional office / civil engineering department.
2. **Build an OECF lookup table** per target device model. Photograph a standard reflectance chart (X-Rite ColorChecker or similar) under controlled illumination; regress pixel value against known luminance; store as a 256-entry table.
3. **Calibrate flash lumen output** per device using a calibrated lux meter at a known distance.
4. **Establish reference targets.** At least 10 physical markings spanning the RL range (freshly painted to badly faded), each measured with the reference instrument.
5. **Fit correction parameters.** Align ReflectScan's output to the reference instrument via a simple linear or piece-wise regression, documented in the repository.

Phase 1 — NHAI pilot on live highway (Month 2–3)

1. **Deployment partner:** one NHAI regional office, one patrol route (~200 km, mixed 4-lane / 6-lane), ~20 patrol phones.
2. **Ground-truth subset:** a sub-sample of ReflectScan segments (e.g. 50 segments) re-measured with the handheld reference instrument within 48 hours of capture.
3. **Metrics.** Report mean absolute error (MAE), bias, and 95% confidence interval between ReflectScan and reference RL per segment. Publish results transparently in the final pilot report.
4. **Failure-mode logging.** Instances where ReflectScan emits UNCAL or its RL disagrees materially from the reference are flagged for post-mortem.
5. **Acceptance criteria for continuing.** The pilot succeeds if ReflectScan identifies the same CRITICAL segments as the reference instrument (binary classification recall > agreed threshold, e.g. > 0.8), even if absolute RL has residual error.

Phase 2 — Integration and broader rollout (Month 4+)

1. **Integration with NHAI Road Asset Management System** so that CRITICAL findings auto-open maintenance tickets.
2. **User training.** Regional engineers learn the dashboard; patrol staff learn the 5-minute app workflow.
3. **Incremental rollout** to additional states, each subject to a repeat of Phase 1-style validation on local reference markings.
4. **Benchmarking against IRC 67 / IRC 35 pass-fail criteria** at each rollout step.

Risk register and mitigations

Risk	Mitigation
Per-device variance exceeds design tolerance	Ship a device allow-list and run the OECF calibration once per model; fall back to UNCAL on unknown models.
Certified instrument unavailable for Phase 0	Partner with CSIR-CRRI, IIT Bombay / Madras transportation labs, or NHAI regional instrumentation stock.
RL accuracy insufficient for absolute compliance reporting	Position ReflectScan as a *screening* tool that flags segments for confirmatory handheld measurement, reducing inspector workload without replacing IRC-compliant sign-off.
Operator misuse (e.g. no mount, handheld shake)	Mount detection check + vibration signature heuristic auto-emits UNCAL on excessively shaky captures.
Regulatory acceptance	Engage the IRC technical committees with pilot data before scaling beyond the pilot state.

What ReflectScan is honestly ready to do *today*

- **Collect dense, GPS-tagged, time-stamped capture data** from any patrol vehicle on any Indian road.
- **Aggregate captures into 100 m segments** and surface them on a map dashboard.
- **Flag captures where conditions invalidate the measurement** via the UNCAL tag, so data consumers never confuse an invalid reading for a valid one.
- **Provide an auditable end-to-end pipeline** on which calibration work can be layered.

What ReflectScan is not ready to do today

- Produce IRC-certified RL values that can be used for contract compliance reporting without corroboration.
- Replace a handheld retroreflectometer in a court-of-law or standards-audit setting.

> **Next step we are looking for:** We would welcome the opportunity to demonstrate ReflectScan alongside an IRC-compliant retroreflectometer at any NHAI site convenient to a regional office. A single calibration session — at a time and location of NHAI's choosing — would let us characterise measurement accuracy and share the findings with NHAI at no cost. No formal commitment is needed from NHAI at this stage; we are simply seeking guidance on whether such a collaboration would be feasible.

OUTPUTS & REPORTING

What the System Produces

Per-measurement payload

```
{
  "lat":         23.0965,
  "lng":         72.5602,
  "rl_value":    87.3,
  "status":      "WARNING",
  "speed_kmh":   62.1,
  "highway":     "NH-48",
  "captured_at": "2026-04-22T17:21:45Z"
}
```

Per-segment payload (backend aggregation)

```
{
  "id":          "seg-a80c...",
  "highway":     "NH-48",
  "rl_avg":      84.2,
  "rl_min":      51.8,
  "rl_max":      118.3,
  "status":      "WARNING",
  "point_count": 38,
  "geometry":    { "type": "LineString", "coordinates": [[...], ...] }
}
```

Classification thresholds

Status	RL range (mcd/m²/lux)	Suggested action
SAFE	> 100	No action; scheduled re-survey
WARNING	54 – 100	Schedule maintenance in the next quarter
CRITICAL	< 54	Resurface; auto-ticket to maintenance system
UNCAL	n/a	Conditions invalid; GPS trace retained; excluded from statistics

Thresholds aligned with *Impact of Road Marking Retroreflectivity on Machine Vision in Dry Conditions* (PMC8963044, 2022) and cross-referenced with IRC 35 / IRC 67. Final thresholds for any deployment would be agreed with NHAI.

Dashboard reporting

1. **Live map** — colour-coded highway segments on OpenStreetMap (Leaflet)
2. **Alert feed** — chronological list of new CRITICAL segments with GPS, timestamp
3. **KPI panel** — % SAFE / WARNING / CRITICAL / UNCAL, coverage %, points captured this week
4. **Export** — CSV (raw segments) and PDF (executive summary per state / division)

Future scope — directions of extension

Direction	How ReflectScan extends
Signboard retroreflectivity (second half of NHAI's problem statement)	Add a `signboard_front` class to the detector. Same physics, same app — one drive covers both markings and signs. Not in the current prototype.
Road studs / delineators	Add a `road_stud` class. Same pipeline, different target geometry.
Predictive maintenance	Once enough time-series data accumulates, regression models on per-segment RL decay can anticipate threshold crossings.
Integration with NHAI ATMS / HTMC	CRITICAL alerts could publish to Variable Message Signs in low-visibility zones.
Citizen contribution mode	Read-only version of the app for civic volunteers; contributions validated before entering the official dataset.
Autonomous-vehicle safety layer	Public API could expose RL scores per lane; AVs could downgrade to manual in low-RL zones.
National safety index	State-wise retroreflectivity index, conceptually comparable to AQI for air quality.

SUMMARY

Why This Solution Wins

ReflectScan is the only submission in this category that combines physics-derived measurement, explicit uncertainty handling, zero-marginal-cost fleet deployment, and an already-live end-to-end system — all in a single Android app that runs on hardware NHAI already operates.

- **Innovation — physics beats black-box ML.** Every comparable smartphone-based approach in published literature uses a neural network to directly regress RL from an image — an unauditable guess. ReflectScan applies the same inverse-square-law photometry a certified retroreflectometer applies, making every intermediate value inspectable and reproducible. The UNCAL tag — a principled refusal to emit a number when conditions are invalid — has no equivalent in any known ML-only proposal.
- **Feasibility — it works today, publicly, verifiably.** The backend API is live (`reflectivity-production.up.railway.app`). The dashboard is live (`reflectivity-3x1c.vercel.app`). The Android APK is installable. 51 automated tests pass. 82 GPS-tagged captures were collected in the field with a two-wheeler stand on 22 April 2026. A judge can independently verify every claim in this document within minutes.
- **Scalability — marginal cost is ~₹1 500 per unit.** One windshield mount per patrol vehicle. No specialist crew. No dedicated vehicle. 200 phones covering one state costs roughly the same as two seats on a single specialized retroreflectometry vehicle. 2 000 phones covering the entire NH network costs ~₹30 Lakh — a fraction of a single year of specialized-vehicle operation.
- **Sustainability — zero new hardware, zero new fuel.** Every capture piggybbacks on a patrol drive NHAI is already running. Open-source stack (Flutter + Python) eliminates vendor lock-in and licence cost. Phones follow standard NHAI replacement cycles.
- **Practical for NHAI — fits the existing operational model.** Patrol drivers need no training beyond a 5-minute walkthrough. Regional engineers see the same NH segments they already manage, now colour-coded by RL status. CRITICAL alerts auto-generate maintenance tickets in RAMS. Nothing in the workflow asks NHAI to change how it operates — it asks only to add one app to existing phones.
- **Honest about limits — no inflated claims.** Rather than claiming IRC-certified accuracy the prototype cannot yet demonstrate, this submission describes exactly what is built, what has been exercised, and what a structured Phase 0 calibration campaign would require. A judge can assess risk with precision — which is itself a mark of engineering maturity.
- **Uniquely positioned as a complement, not a replacement.** ReflectScan does not ask NHAI to retire its certified retroreflectometers. It proposes to screen the entire network densely and cheaply, so that certified instruments are deployed only where the data says they are needed — multiplying their effective reach without additional capex.

> **One-line pitch:** ReflectScan turns every NHAI patrol vehicle into a retroreflectivity sensor at a marginal cost of ₹1 500 — using physics, not guesswork, with a live system a judge can open right now.

VERIFICATION & CONTACT

How to Verify Every Claim in This Document

Everything in this submission can be independently checked by a reviewer.

Live resources

Resource	Link
Public GitHub repository (full source)	`https://github.com/kuntardivyang/reflectivity`
Screenshots, videos & APK (Google Drive)	`https://drive.google.com/drive/folders/1xpssEgShkJ3bRBdnN1gJgEUiT1i-ei5d`
Live dashboard	`https://reflectivity-3x1c.vercel.app`
Live API + OpenAPI docs	`https://reflectivity-production.up.railway.app/docs`
Android APK (direct install)	`Google Drive folder → reflectscan.apk`
Architecture document	`/docs/architecture.md in repo`
Research citations	`/docs/research_findings.md in repo`

Reproducibility

```
# Backend (requires Postgres + PostGIS)
cd backend
python3 -m venv .venv && source .venv/bin/activate
pip install -r requirements.txt
python -m db.init_db
python -m scripts.seed_nh48
uvicorn main:app --reload

# Mobile (requires Flutter 3.19+)
cd mobile
flutter pub get
flutter build apk --release
# install resulting app-release.apk on any Android 8+ device

# Tests
cd backend && pytest          # 27 passing in ~1 s
cd ../mobile && flutter test  # 24 passing in ~1 s
```

Evidence artefacts packaged with this submission

- `01_Proposal.pdf` — this document
- `03_Prototype_Screenshots/` — DEMO MODE banner, live survey HUD, session history, session detail, UNCAL status
- `04_Pipeline_Validation_Ride/` — metadata for the 82-capture session collected from a two-wheeler with mobile-phone stand on 22 April 2026, Ahmedabad local roads (pipeline-validation only; see page 7 for scope)
- `05_Architecture_Diagram.png` — high-resolution export of the Layer 1 / 2 / 3 diagram
- `06_Cost_Comparison_Chart.png` — 5-year projected cost trajectory
- `README.txt` — one-page index with every link

Evaluation-criteria map

NHAI criterion	Where addressed
Innovation	Page 5 — four design contributions, head-to-head positioning
Feasibility	Pages 6 & 7 — component status, pipeline-validation walk, validation scope, deployment roadmap
Scalability & Sustainability	Page 8 — projected scale and cost trajectory, data re-use, sustainability
Presentation & Documentation	This document + repo docs + OpenAPI spec + screenshots + field-walk metadata

Contact

Team Leader: Divyang Kuntar
Email: divyangkuntar@gmail.com
Mobile: 7016509214
Project repository: https://github.com/kuntardivyang/reflectivity

Thank you for reviewing ReflectScan.
The prototype is the output of four days of design, implementation, and pipeline validation on local road conditions (two-wheeler with mobile-phone stand). It is offered as an honest proof-of-concept. The numbers are the numbers; the open source is the open source; the calibration work is the calibration work. We welcome the opportunity to do the Phase 0 calibration and Phase 1 pilot in partnership with NHAI.